

Closest Pair of Points

(from “Algorithm Design” by J. Kleinberg and E. Tardos)

Closest pair. Given n points in the plane, find a pair with smallest Euclidean distance between them.

Fundamental geometric primitive.

- Graphics, computer vision, geographic information systems, molecular modeling, air traffic control.
- Special case of nearest neighbor, Euclidean MST, Voronoi.

Brute force. Check all pairs of points p and q with $\Theta(n^2)$ comparisons.

1-D version. $O(n \log n)$ easy if points are on a line.


Assumption. No two points have same x coordinate.

↑
to make presentation cleaner

Closest Pair of Points

Algorithm.


- *Divide:* draw vertical line \mathcal{L} so that roughly $\frac{1}{2}n$ points on each side.



Closest Pair of Points

Algorithm.

- *Divide*: draw vertical line \mathcal{L} so that roughly $\frac{1}{2}n$ points on each side.
- *Conquer*: find closest pair in each side recursively.




Closest Pair of Points

Algorithm.


- *Divide*: draw vertical line \mathcal{L} so that roughly $\frac{1}{2}n$ points on each side.
- *Conquer*: find closest pair in each side recursively.
- *Combine*: find closest pair with one point in each side.
- Return best of 3 solutions.

← seems like $\Theta(n^2)$



Closest Pair of Points


Find closest pair with one point in each side, *assuming that distance* $< \delta$.



Closest Pair of Points

Find closest pair with one point in each side, *assuming that distance* $< \delta$.


- Observation: only need to consider points within δ of line \mathcal{L} .



Closest Pair of Points

Find closest pair with one point in each side, *assuming that distance* $< \delta$.


- Observation: only need to consider points within δ of line \mathcal{L} .
- Sort points in 2δ -strip by their y coordinate.



Closest Pair of Points

Find closest pair with one point in each side, *assuming that distance $< \delta$* .

- Observation: only need to consider points within δ of line \mathcal{L} .
- Sort points in 2δ -strip by their y coordinate.
- Only check distances of those within 11 positions in sorted list!



Closest Pair of Points


Def. Let s_i be the point in the 2δ -strip, with the i^{th} smallest y -coordinate.

Claim. If $|i - j| \geq 12$, then the distance between s_i and s_j is at least δ .

Pf.

- No two points lie in same $\frac{1}{2}\delta$ -by- $\frac{1}{2}\delta$ box.
- Two points at least 2 rows apart have distance $\geq 2(\frac{1}{2}\delta)$. ■

Fact. Still true if we replace 12 with 7.



Closest Pair Algorithm

Closest-Pair(p_1, \dots, p_n) {

Compute separation line L such that half the points are on one side and half on the other side.

$O(n \log n)$

$\delta_1 = \text{Closest-Pair}(\text{left half})$

$2T(n/2)$

$\delta_2 = \text{Closest-Pair}(\text{right half})$

$\delta = \min(\delta_1, \delta_2)$

Delete all points further than δ from separation line L

$O(n)$

Sort remaining points by y-coordinate.

$O(n \log n)$

Scan points in y-order and compare distance between each point and next 11 neighbors. If any of these distances is less than δ , update δ .

$O(n)$

return δ .

}

Closest Pair of Points: Analysis

Running time.

$$T(n) \leq 2T(n/2) + O(n \log n) \Rightarrow T(n) = O(n \log^2 n)$$

Q. Can we achieve $O(n \log n)$?

A. Yes. Don't sort points in strip from scratch each time.

- *Each recursive returns all points sorted by y coordinate.*
- *Sort by **merging** two pre-sorted lists.*

$$T(n) \leq 2T(n/2) + O(n) \Rightarrow T(n) = O(n \log n)$$